

Effective deep learning-based multi-modal retrieval

Wei Wang¹ · Xiaoyan Yang² · Beng Chin Ooi¹ ·
Dongxiang Zhang¹ · Yueting Zhuang³

Received: 1 January 2015 / Revised: 19 May 2015 / Accepted: 15 June 2015 / Published online: 19 July 2015
© Springer-Verlag Berlin Heidelberg 2015

Abstract Multi-modal retrieval is emerging as a new search paradigm that enables seamless information retrieval from various types of media. For example, users can simply snap a movie poster to search for relevant reviews and trailers. The mainstream solution to the problem is to learn a set of mapping functions that project data from different modalities into a common metric space in which conventional indexing schemes for high-dimensional space can be applied. Since the effectiveness of the mapping functions plays an essential role in improving search quality, in this paper, we exploit deep learning techniques to learn effective mapping functions. In particular, we first propose a general learning objective that effectively captures both intramodal and intermodal semantic relationships of data from heterogeneous sources. Given the general objective, we propose two learning algorithms to realize it: (1) an unsupervised approach that uses stacked auto-encoders and requires minimum prior knowledge on the training data and (2) a

supervised approach using deep convolutional neural network and neural language model. Our training algorithms are memory efficient with respect to the data volume. Given a large training dataset, we split it into mini-batches and adjust the mapping functions continuously for each batch. Experimental results on three real datasets demonstrate that our proposed methods achieve significant improvement in search accuracy over the state-of-the-art solutions.

Keywords Deep learning · Multi-modal retrieval · Hashing · Auto-encoders · Deep convolutional neural network · Neural language model

1 Introduction

The prevalence of social networking has significantly increased the volume and velocity of information shared on the Internet. A tremendous amount of data in various media types is being generated every day in social networking systems, and images and video contribute the main bulk of the data. For instance, Twitter recently reported that over 340 million tweets were sent each day,¹ while Facebook reported that around 300 million photographs were created each day.² These data, together with other domain specific data, such as medical data, surveillance and sensory data, are big data that can be exploited for insights and contextual observations. However, effective retrieval of such huge amounts of media from heterogeneous sources remains a big challenge.

✉ Wei Wang
wangwei@comp.nus.edu.sg

Xiaoyan Yang
xiaoyan.yang@adsc.com.sg

Beng Chin Ooi
ooibc@comp.nus.edu.sg

Dongxiang Zhang
zhangdo@comp.nus.edu.sg

Yueting Zhuang
yzhuang@zju.edu.cn

¹ School of Computing, National University of Singapore, Singapore, Singapore

² Advanced Digital Sciences Center, Illinois at Singapore Pte, Singapore, Singapore

³ College of Computer Science and Technology, Zhejiang University, Hangzhou, China

¹ <https://blog.twitter.com/2012/twitter-turns-six>.

² <http://techcrunch.com/2012/08/22/how-big-is-facebooks-data-2-5-billion-pieces-of-content-and-500-terabytes-ingested-every-day/>.

In this paper, we exploit deep learning techniques, which have been successfully applied in processing media data [3, 29, 33], to solve the problem of large-scale information retrieval from multiple modalities. Each modality represents one type of media such as text, image or video. Depending on the heterogeneity of data sources, we have two types of searches:

1. **Intramodal search** has been extensively studied and widely used in commercial systems. Examples include web document retrieval via keyword queries and content-based image retrieval.
2. **Cross-modal search** enables users to explore relevant resources from different modalities. For example, a user can use a tweet to retrieve relevant photographs and videos from other heterogeneous data sources, or search relevant textual descriptions or videos by submitting an interesting image as a query.

There has been a long stream of research on multi-modal retrieval [4, 21, 25, 30, 36, 43–45]. These works follow the same query processing strategy, which consists of two major steps. First, a set of mapping functions are learned to project data from **different modalities** into a **common latent space**. Second, a multi-dimensional index for each modality in the common space is built for efficient similarity retrieval. Since the second step is a classic k NN problem and has been extensively studied [16, 40, 42], we focus on the optimization of the first step and propose two types of novel mapping functions based on deep learning techniques.

We propose a general learning objective that effectively captures both intramodal and intermodal semantic relationships of data from heterogeneous sources. In particular, we differentiate modalities in terms of their representations' ability to capture semantic information and robustness when noisy data are involved. The modalities with better representations are assigned with larger weight for the sake of learning more effective mapping functions. Based on the objective function, we design an unsupervised algorithm using stacked auto-encoders (SAEs). SAE is a deep learning model that has been widely applied in many unsupervised feature learning and classification tasks [13, 31, 34, 38]. If the media are annotated with semantic labels, we design a supervised algorithm to realize the learning objective. The supervised approach uses a deep convolutional neural network (DCNN) and neural language model (NLM). It exploits the label information, thus can learn robust mapping functions against noisy input data. DCNN and NLM have shown great success in learning image features [8, 10, 20] and text features [28, 33], respectively.

Compared with existing solutions for multi-modal retrieval, our approaches exhibit three major advantages. First, our mapping functions are nonlinear, which are more expressive

than the linear projections used in IMH [36] and CVH [21]. The deep structures of our models can capture more abstract concepts at higher layers, which is very useful in modeling categorical information of data for effective retrieval. Second, we require minimum prior knowledge in the training. Our unsupervised approach only needs relevant data pairs from different modalities as the training input. The supervised approach requires additional labels for the media objects. In contrast, MLBE [43] and IMH [36] require a big similarity matrix of intramodal data for each modality. LSCMR [25] uses training examples, each of which consists of a list of objects ranked according to their relevance (based on manual labels) to the first one. Third, our training process is memory efficient because we split the training dataset into mini-batches and iteratively load and train each mini-batch in memory. However, many existing works (e.g., CVH, IMH) have to load the whole training dataset into memory which is infeasible when the training dataset is too large.

In summary, main contributions of this paper are:

- We propose a general learning objective for learning mapping functions to project data from different modalities into a common latent space for multi-modal retrieval. The learning objective differentiates modalities in terms of their input features' quality of capturing semantics.
- We realize the general learning objective by one unsupervised approach and one supervised approach based on deep learning techniques.
- We conduct extensive experiments on three real datasets to evaluate our proposed mapping mechanisms. Experimental results show that the performance of our method is superior to state-of-the-art methods.

The remainder of the paper is organized as follows. Problem statements and overview are provided in Sects. 2 and 3. After that, we describe the unsupervised and supervised approaches in Sects. 4 and 5, respectively. Query processing is presented in Sect. 6. We discuss related works in Sect. 7 and present our experimental study in Sect. 8. We conclude our paper in Sect. 9. This work is an extended version of [39] which proposed an unsupervised learning algorithm for multi-modal retrieval. We add a supervised learning algorithm as a new contribution in this work. We also revise Sect. 3 to unify the learning objective of the two approaches. Sections 5 and 8.3 are newly added for describing the supervised approach and its experimental results.

2 Problem statements

In our data model, the database \mathbb{D} consists of objects from multiple modalities. For ease of presentation, we use images and text as two sample modalities to explain our idea, i.e., we

assume that $\mathbb{D} = \mathbb{D}_I \cup \mathbb{D}_T$. An image (resp. a text document) is represented by a feature vector $x \in \mathbb{D}_I$ (resp. $y \in \mathbb{D}_T$). To conduct multi-modal retrieval, we need a relevance measurement for the query and the database object. However, the database consists of objects from different modalities, there is no such widely accepted measurement. A common approach is to learn a set of mapping functions that project the original feature vectors into a common latent space such that semantically relevant objects (e.g., image and its tags) are located close. Consequently, our problem includes the following two sub-problems.

Definition 1 Common Latent Space Mapping

Given an image $x \in \mathbb{D}_I$ and a text document $y \in \mathbb{D}_T$, find two mapping functions $f_I : \mathbb{D}_I \rightarrow \mathbb{Z}$, and $f_T : \mathbb{D}_T \rightarrow \mathbb{Z}$, such that if x and y are semantically relevant, the distance between $f_I(x)$ and $f_T(y)$ in the common latent space \mathbb{Z} , denoted by $\text{dist}_{\mathbb{Z}}(f_I(x), f_T(y))$, is small.

The common latent space mapping provides a unified approach to measuring distance of objects from different modalities. As long as all objects can be mapped into the same latent space, they become comparable. Once the mapping functions f_I and f_T have been determined, the multi-modal search can then be transformed into the classic k NN problem, defined as follows

Definition 2 Multi-Modal Search

Given a query object $Q \in \mathbb{D}_q$ and a target domain \mathbb{D}_t ($q, t \in \{I, T\}$), find a set $O \subset \mathbb{D}_t$ with k objects such that $\forall o \in O$ and $o' \in \mathbb{D}_t/O$, $\text{dist}_{\mathbb{Z}}(f_q(Q), f_t(o')) \geq \text{dist}_{\mathbb{Z}}(f_q(Q), f_t(o))$.

Since both q and t have two choices, four types of queries can be derived, namely $Q_{q \rightarrow t}$ and $q, t \in \{I, T\}$. For instance, $Q_{I \rightarrow T}$ searches relevant text in \mathbb{D}_T given an image from \mathbb{D}_I . By mapping objects from different high-dimensional feature spaces into a low-dimensional latent space, queries can be efficiently processed using existing multi-dimensional indexes [16, 40]. Our goal is then to learn a set of effective mapping functions which preserve well both intramodal semantics (i.e., semantic relationships within each modality) and intermodal semantics (i.e., semantic relationships across modalities) in the latent space. The effectiveness of mapping functions is measured by the accuracy of multi-modal retrieval using latent features.

3 Overview of multi-modal retrieval

The flowchart of our multi-modal retrieval framework is illustrated in Fig. 1. It consists of three main steps: (1) offline model training (2) offline indexing and (3) online k NN query processing. In step 1, relevant image-text pairs are used as

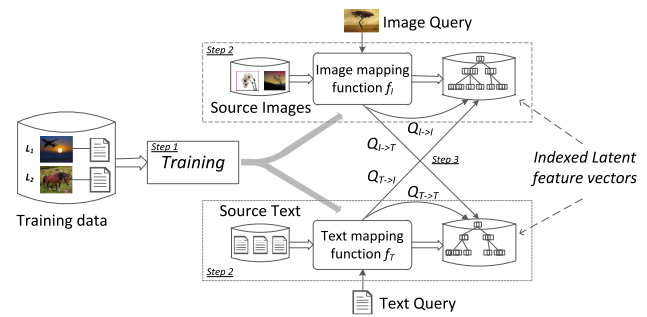


Fig. 1 Flowchart of multi-modal retrieval framework. Step 1 is offline model training that learns mapping functions. Step 2 is offline indexing that maps source objects into latent features and creates proper indexes. Step 3 is online multi-modal k NN query processing

input training data for learning the mapping functions. For example, image-text pairs can be collected from Flickr where the text features are extracted from tags and descriptions for images. If they are associated with additional semantic labels (e.g., categories), we use a supervised training algorithm. Otherwise, an unsupervised training algorithm is used. After step 1, we can obtain a mapping function $f_m : \mathbb{D}_m \rightarrow \mathbb{Z}$ for each modality $m \in \{I, T\}$. In step 2, objects from different modalities are first mapped into the common space \mathbb{Z} by function f_m . With such unified representation, the latent features from the same modality are then inserted into a high-dimensional index for k NN query processing. When a query $Q \in \mathbb{D}_m$ comes, it is first mapped into \mathbb{Z} using its modal-specific mapping function f_m . Based on the query type, k nearest neighbors are retrieved from the index built for the target modality and returned to the user. For example, image index is used for queries of type $Q_{I \rightarrow I}$ and $Q_{T \rightarrow I}$ against the image database.

General learning objective A good objective function plays a crucial role in learning effective mapping functions. In our multi-modal search framework, we design a general learning objective function \mathcal{L} . By taking into account the image and text modalities, our objective function is defined as follows:

$$\mathcal{L} = \beta_I \mathcal{L}_I + \beta_T \mathcal{L}_T + \mathcal{L}_{I,T} + \xi(\theta) \quad (1)$$

where $\mathcal{L}_m, m \in \{I, T\}$ is called the intramodal loss to reflect how well the intramodal semantics are captured by the latent features. The smaller the loss, the more effective the learned mapping functions are, $\mathcal{L}_{I,T}$ is called the intermodal loss which is designed to capture intermodal semantics. The last term is used as regularization to prevent over-fitting [14] (L_2 norm is used in our experiment). θ denotes all parameters involved in the mapping functions. $\beta_m, m \in \{I, T\}$ denotes the weight of the loss for modality m in the objective function. We observe in our training process that assigning different weights to different modalities according to the nature of its

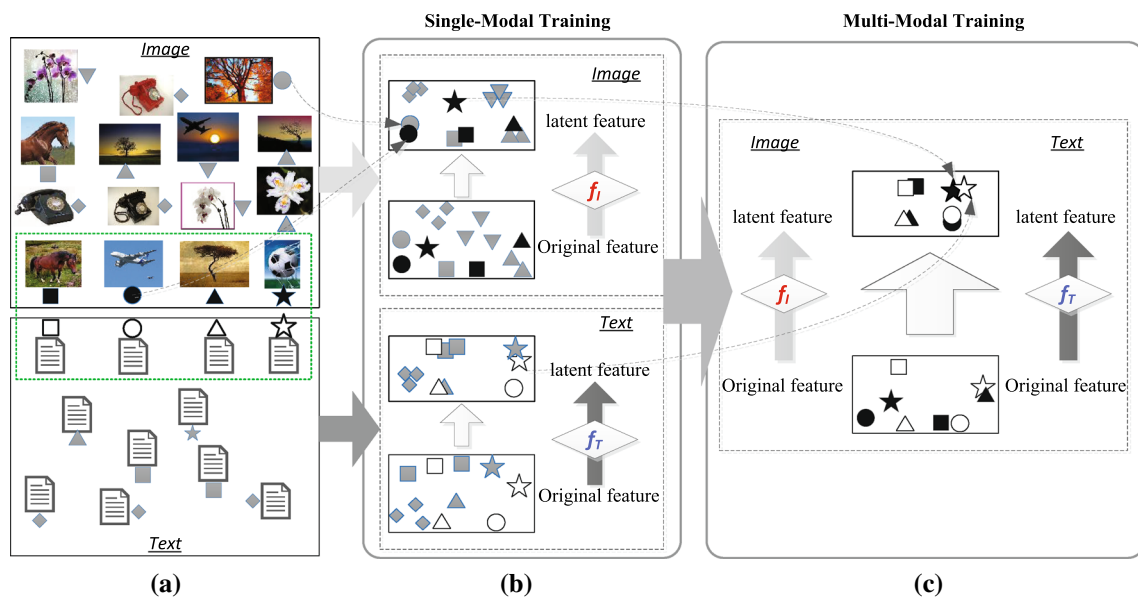


Fig. 2 Flowchart of training. Relevant images (or text) are associated with the same shape (e.g., ■). In single-modal training, objects of same shape and modality are moving close to each other. In multi-modal

training, objects of same shape from all modalities are moving close to each other. **a** Input. **b** Training stage I. **c** Training stage II

data offers better performance than treating them equally. For the modality with lower-quality input feature (due to noisy data or poor data representation), we assign smaller weight for its intramodal loss in the objective function. The intuition of setting β_I and β_T in this way is that, by relaxing the constraints on intramodal loss, we enforce the intermodal constraints. Consequently, the intramodal semantics of the modality with lower-quality input feature can be preserved or even enhanced through their intermodal relationships with high-quality modalities. Details of setting β_I and β_T will be discussed in Sects. 4.3 and 5.3.

Training Training is to find the optimal parameters involved in the mapping functions that minimizes \mathcal{L} . Two types of mapping functions are proposed in this paper. One is trained by an unsupervised algorithm, which uses simple image–text pairs for training. No other prior knowledge is required. The other one is trained by a supervised algorithm which exploits additional label information to learn robust mapping functions against noisy training data. For both mapping functions, we design a two-stage training procedure to find the optimal parameters. A complete training process is illustrated in Fig. 2. In *stage I*, one mapping function is trained independently for each modality with the objective to map similar features in one modality close to each other in the latent space. This training stage serves as the pre-training of stage II by providing a good initialization for the parameters. In *stage II*, we jointly optimize Eq. 1 to capture both intramodal semantics and intermodal semantics. The learned mapping functions project semantically relevant objects close to each other in the latent space as shown in the figure.




Fig. 3 Model of MSAE, which consists of one SAE for each modality. The trained SAE maps input data into latent features

4 Unsupervised approach: MSAE

In this section, we propose an unsupervised learning algorithm called multi-modal stacked auto-encoders (MSAE) to learn the mapping function f_I and f_T . The model is shown in Fig. 3. We first present the preliminary knowledge of auto-encoder and stacked auto-encoders. Based on stacked auto-encoders, we address how to define the terms \mathcal{L}_I , \mathcal{L}_T and $\mathcal{L}_{I,T}$ in our general objective learning function in Eq. 1.

4.1 Background: auto-encoder and stacked auto-encoder

Auto-encoder Auto-encoder has been widely used in unsupervised feature learning and classification tasks [13, 31, 34, 38]. It can be seen as a special neural network with three layers—the input layer, the latent layer and the reconstruction layer. As shown in Fig. 4, the raw input feature $x_0 \in \mathcal{R}^{d_0}$ in the input layer is **encoded** into latent feature $x_1 \in \mathcal{R}^{d_1}$ via




Fig. 4 Auto-encoder

a deterministic mapping f_e :

$$x_1 = f_e(x_0) = s_e(W_1^T x_0 + b_1) \quad (2)$$

where s_e is the activation function of the encoder, $W_1 \in \mathbb{R}^{d_0 \times d_1}$ is a weight matrix, and $b_1 \in \mathcal{R}^{d_1}$ is a bias vector. The latent feature x_1 is then **decoded** back to $x_2 \in \mathcal{R}^{d_0}$ via another mapping function f_d :

$$x_2 = f_d(x_1) = s_d(W_2^T x_1 + b_2) \quad (3)$$

Similarly, s_d is the activation function of the decoder with parameters $\{W_2, b_2\}$, $W_2 \in \mathbb{R}^{d_1 \times d_0}$, $b_2 \in \mathcal{R}^{d_0}$. Sigmoid function or Tanh function is typically used as the activation functions s_e and s_d . The parameters $\{W_1, W_2, b_1, b_2\}$ of the auto-encoder are learned with the objective of minimizing the difference (called reconstruction error) between the raw input x_0 and the reconstruction output x_2 . Squared Euclidean distance, negative log likelihood and cross-entropy are often used to measure the reconstruction error. By minimizing the reconstruction error, we can use the latent feature to reconstruct the original input with minimum information loss. In this way, the latent feature preserves regularities (or semantics) of the input data.

Stacked auto-encoder Stacked auto-encoders (SAE) are constructed by stacking multiple (e.g., h) auto-encoders. The input feature vector x_0 is fed to the bottom auto-encoder. After training the bottom auto-encoder, the latent representation x_1 is propagated to the higher auto-encoder. The same procedure is repeated until all the auto-encoders are trained. The latent representation x_h from the top (i.e., h -th) auto-encoder is the output of the stacked auto-encoders, which can be further fed into other applications, such as SVM for classification. The stacked auto-encoders can be fine-tuned by minimizing the reconstruction error between the input feature x_0 and the reconstruction feature x_{2h} which is computed by forwarding the x_0 through all encoders and then through all decoders as shown in Fig. 5. In this way, the output feature x_h can reconstruct the input feature with minimal information loss. In other words, x_h preserves regularities (or semantics) of the input data x_0 .




Fig. 5 Fine-tune stacked auto-encoders


4.2 Realization of the learning objective in MSAE

4.2.1 Modeling intramodal semantics of data

We extend SAEs to model the intramodal losses in the general learning objective (Eq. 1). Specifically, \mathcal{L}_I and \mathcal{L}_T are modeled as the reconstruction errors for the image SAE and the text SAE, respectively. Intuitively, if the two reconstruction errors are small, the latent features generated by the top auto-encoder would be able to reconstruct the original input well and, consequently, capture the regularities of the input data well. This implies that with small reconstruction error, two objects from the same modality that are similar in the original space would also be close in the latent space. In this way, we are able to capture the intramodal semantics of data by minimizing \mathcal{L}_I and \mathcal{L}_T , respectively. But to use the SAEs we have to design the decoders of the bottom auto-encoders carefully to handle different input features.

The raw (input) feature of an image is a high-dimensional real-valued vector (e.g., color histogram or bag-of-visual-words). In the encoder, each input image feature is mapped to a latent vector using Sigmoid function as the activation function s_e (Eq. 2). However, in the decoder, the Sigmoid activation function, whose range is $[0,1]$, performs poorly on reconstruction because the raw input unit (referring to one dimension) is not necessarily within $[0,1]$. To solve this issue, we follow Hinton [14] and model the raw input unit as a linear unit with independent Gaussian noise. As shown in Fig. 6a, the average unit value of image feature typically follows Gaussian distribution. When the input data are normalized with zero mean and unit variance, the Gaussian noise term can be omitted. In this case, we use an identity function for the activation function s_d in the bottom decoder. Let x_0 denote the input image feature vector, and let x_{2h} denote the feature vector reconstructed from the top latent feature x_h (h is the depth of the stacked auto-encoders). Using Euclidean distance to measure the reconstruction error, we define \mathcal{L}_I for x_0 as:

Fig. 6 Distribution of image (a) and text (b) features extracted from NUS-WIDE training dataset (See Sect. 8). Each figure is generated by averaging the units for each feature vector, and then the histogram for all data is plotted



$$\mathcal{L}_I(x_0) = \|x_0 - x_{2h}\|_2^2 \quad (4)$$

The raw (input) feature of text is a word count vector or tag occurrence vector.³ We adopt the Rate Adapting Poisson model [32] for reconstruction because the histogram for the average value of text input unit generally follows Poisson distribution (Fig. 6b). In this model, the activation function in the bottom decoder is

$$x_{2h} = s_d(z_{2h}) = l \frac{e^{z_{2h}}}{\sum_j e^{z_{2h_j}}} \quad (5)$$

where $l = \sum_j x_{0j}$ is the number of words in the input text and $z_{2h} = W_{2h}^T x_{2h-1} + b_{2h}$. The probability of a reconstruction unit x_{2h_i} being the same as the input unit x_{0_i} is:

$$p(x_{2h_i} = x_{0_i}) = \text{Pois}(x_{0_i}, x_{2h_i}) \quad (6)$$

where $\text{Pois}(n, \lambda) = \frac{e^{-\lambda} \lambda^n}{n!}$. Based on Eq. 6, we define \mathcal{L}_T using negative log likelihood:

$$\mathcal{L}_T(x_0) = -\log \prod_i p(x_{2h_i} = x_{0_i}) \quad (7)$$

By minimizing \mathcal{L}_T , we require x_{2h} to be similar as x_0 . In other words, the latent feature x_h is trained to reconstruct the input feature well and thus preserves the regularities of the input data well.

4.2.2 Modeling intermodal semantics of data

Given one relevant image–text pair (x_0, y_0) , we forward them through the encoders of their stacked auto-encoders to generate latent feature vectors (x_h, y_h) (h is the height of the SAE). The intermodal loss is then defined as,

³ The binary value for each dimension indicates whether the corresponding tag appears or not.

$$\mathcal{L}_{I,T}(x_0, y_0) = \text{dist}(x_h, y_h) = \|x_h - y_h\|_2^2 \quad (8)$$

By minimizing $\mathcal{L}_{I,T}$, we capture the intermodal semantics of data. The intuition is quite straightforward: If two objects x_0 and y_0 are relevant, the distance between their latent features x_h and y_h shall be small.

4.3 Training

Following the training flow shown in Fig. 2, in stage I we train a SAE for the image modality and a SAE for the text modality separately. Back-propagation [22] (see “Appendix”) is used to calculate the gradients of the objective loss, i.e., \mathcal{L}_I or \mathcal{L}_T , w.r.t. the parameters. Then the parameters are updated according to mini-batch stochastic gradient descent (SGD) (see “Appendix”), which averages the gradients contributed by a mini-batch of training records (images or text documents) and then adjusts the parameters. The learned image and text SAEs are fine-tuned in stage II by back-propagation and mini-batch SGD with the objective to find the optimal parameters that minimize the learning objective (Eq. 1). In our experiment, we observe that the training would be more stable if we alternatively adjust one SAE with the other SAE fixed.

Setting β_I & β_T β_I and β_T are the weights of the reconstruction error of image and text SAEs, respectively, in the objective function (Eq. 1). As mentioned in Sect. 3, they are set based on the quality of each modality’s raw (input) feature. We use an example to illustrate the intuition. Consider a relevant object pair (x_0, y_0) from modality x and y . Assume x ’s feature is of low quality in capturing semantics (e.g., due to noise), while y ’s feature is of high quality. If x_h and y_h are the latent features generated by minimizing the reconstruction error, then y_h can preserve the semantics well, while x_h is not as meaningful due to the low quality of x_0 . To solve this problem, we combine the intermodal distance between x_h and y_h in the learning objective function and assign smaller weight to the reconstruction error of x_0 . This is the same as

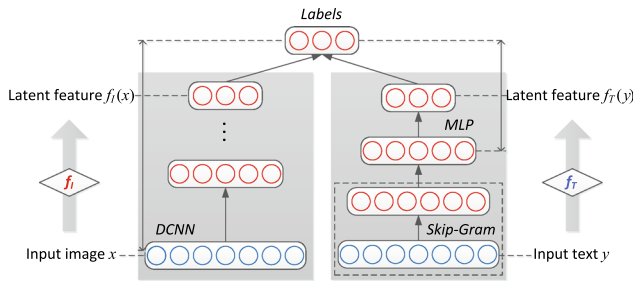


Fig. 7 Model of MDNN, which consists of one DCNN for image modality, and one skip-gram + MLP for text modality. The trained DCNN (or skip-gram + MLP) maps input data into latent features

increasing the weight of the intermodal distance from x_h to y_h . As a result, the training algorithm will move x_h toward y_h to make their distance smaller. In this way, the semantics of low quality x_h could be enhanced by the high-quality feature y_h .

In the experiment, we evaluate the quality of each modality's raw feature on a validation dataset by performing intramodal search against the latent features learned in single-modal training. Modality with worse search performance is assigned a smaller weight. Notice that because the dimensions of the latent space and the original space are usually of different orders of magnitude, the scale of \mathcal{L}_I , \mathcal{L}_T and $\mathcal{L}_{I,T}$ are different. In the experiment, we also scale β_I and β_T to make the losses comparable, i.e., within an order of magnitude.

5 Supervised approach: MDNN

In this section, we propose a supervised learning algorithm called multi-modal deep neural network (MDNN) based on a deep convolutional neural network (DCNN) model and a neural language model (NLM) to learn mapping functions for the image modality and the text modality, respectively. The model is shown in Fig. 7. First, we provide some background on DCNN and NLM. Second, we extend one DCNN [20] and one NLM [28] to model intramodal losses involved in the general learning objective (Eq. 1). Third, the intermodal loss is specified and combined with the intramodal losses to realize the general learning objective. Finally, we describe the training details.

5.1 Background: deep convolutional neural network and neural language model

Deep convolutional neural network (DCNN) DCNN has shown great success in computer vision tasks [8, 10] since the first DCNN (called AlexNet) was proposed by Alex [20]. It has specialized connectivity structure, which usually consists

of multiple convolutional layers followed by fully connected layers. These layers form stacked, multiple-staged feature extractors, with higher layers generating more abstract features from lower ones. On top of the feature extractor layers, there is a classification layer. Please refer to [20] for a more comprehensive review of DCNN.

The input to DCNN is raw image pixels such as an RGB vector, which is forwarded through all feature extractor layers to generate a feature vector that is a high-level abstraction of the input data. The training data of DCNN consist of image-label pairs. Let x denote the image raw feature and $f_I(x)$ the feature vector extracted from DCNN. t is the binary label vector of x . If x is associated with the i -th label l_i , t_i is set to 1 and all other elements are set to 0. $f_I(x)$ is forwarded to the classification layer to predict the final output $p(x)$, where $p_i(x)$ is the probability of x being labeled with l_i . Given x and $f_I(x)$, $p_i(x)$ is defined as:

$$p_i(x) = \frac{e^{f_I(x)_i}}{\sum_j e^{f_I(x)_j}} \quad (9)$$

which is a softmax function. Based on Eq. 9, we define the prediction error, or *softmax loss* as the negative log likelihood:

$$\mathcal{L}_I(x, t) = - \sum_i t_i \log p_i(x) \quad (10)$$

Neural Language Model (NLM) NLMs, first introduced in [21], learn a dense feature vector for each word or phrase, called a distributed representation or a *word embedding*. Among them, the skip-gram model (SGM) [28] proposed by Mikolov et al. is the state-of-the-art. Given two words a and b that co-occur, SGM models the conditional probability $p(a|b)$ using softmax:

$$p(a|b) = \frac{e^{v_a \cdot v_b}}{\sum_{\tilde{a}} e^{v_{\tilde{a}} \cdot v_b}} \quad (11)$$

where v_a and v_b are vector representations of a word a and a context b , respectively. The denominator $\sum_{\tilde{a}} e^{v_{\tilde{a}} \cdot v_b}$ is expensive to calculate given a large vocabulary, where \tilde{a} is any word in the vocabulary. Thus, approximations were proposed to estimate it [8]. Given a corpus of sentences, SGM is trained to learn vector representations v by maximizing Eq. 11 over all co-occurring pairs.

The learned dense vectors can be used to construct a dense vector for one sentence or document (e.g., by averaging) or to calculate the similarity of two words, e.g., using the cosine similarity function.

5.2 Realization of the learning objective in MDNN

5.2.1 Modeling intramodal semantics of data

Having witnessed the outstanding performance of DCNNs in learning features for visual data [8, 10], and NLMs in learning features for text data [33], we extend one instance of DCNN—AlexNet [20] and one instance of NLM—Skip-Gram model (SGM) [28] to model the intramodal semantics of images and text, respectively.

Image We employ AlexNet to serve as the mapping function f_I for image modality. An image x is represented by an RGB vector. The feature vector $f_I(x)$ learned by AlexNet is used to predict the associated labels of x . However, the objective of the original AlexNet is to predict single label of an image, while in our case images are annotated with multiple labels. We thus follow [11] to extend the softmax loss (Eq. 10) to handle multiple labels as follows:

$$\mathcal{L}_I(x, t) = -\frac{1}{\sum_i t_i} \sum_i t_i \log p_i(x) \quad (12)$$

where $p_i(x)$ is defined in Eq. 9. Different from SAE, which models reconstruction error to preserve intramodal semantics, the extended AlexNet tries to minimize the prediction error \mathcal{L}_I shown in Eq. 12. By minimizing prediction error, we require the learned high-level feature vectors $f_I(x)$ to be discriminative in predicting labels. Images with similar labels shall have similar feature vectors. In this way, the intramodal semantics are preserved.

Text We extend SGM to learn the mapping function f_T for text modality. Due to the noisy nature of text (e.g., tags) associated with images [23], directly training the SGM over the tags would carry noise into the learned features. However, labels associated with images are carefully annotated by humans and thus are more accurate. Hence, we extend the SGM to integrate label information so as to learn robust features against noisy text (tags). The main idea is, we first train a SGM [28], treating all tags associated with one image as an input sentence. After training, we obtain one word embedding for each tag. By averaging word embeddings of all tags of one image, we create one text feature vector for those tags. Second, we build a Multi-layer perceptron (MLP) with two hidden layers on top of the SGM. The text feature vectors are fed into the MLP to predict image labels. Let y denote the input text (e.g., a set of image tags), and let \tilde{y} denote the averaged word embedding generated by SGM for tags in y . MLP together with SGM serves as the mapping function f_T for the text modality,

$$f_T(y) = W_2 \cdot s(W_1 \tilde{y} + b_1) + b_2 \quad (13)$$

$$s(v) = \max(0, v) \quad (14)$$

where W_1 and W_2 are weight matrices, b_1 and b_2 are bias vectors, and $s()$ is the ReLU activation function [20].⁴ The loss function of MLP is similar to that of the extended AlexNet for image-label prediction:

$$\mathcal{L}_T(y, t) = -\frac{1}{\sum_i t_i} \sum_i \log q_i(y) \quad (15)$$

$$q_i(y) = \frac{e^{f_T(y)_i}}{\sum_j e^{f_T(y)_j}} \quad (16)$$

We require the learned text latent features $f_T(y)$ to be discriminative for predicting labels. In this way, we model the intramodal semantics for the text modality.⁵

5.2.2 Modeling intermodal semantics

After extending the AlexNet and skip-gram model to preserve the intramodal semantics for images and text, respectively, we jointly learn the latent features for image and text to preserve the intermodal semantics. We follow the general learning objective in Eq. 1 and realize \mathcal{L}_I and \mathcal{L}_T using Eqs. 12 and 15, respectively. Euclidean distance is used to measure the difference of the latent features for an image–text pair, i.e., $\mathcal{L}_{I,T}$ is defined similarly as in Eq. 8. By minimizing the distance of latent features for an image–text pair, we require their latent features to be closer in the latent space. In this way, the intermodal semantics are preserved.

5.3 Training

Similar to the training of MSAE, the training of MDNN consists of two steps. The first step trains the extended AlexNet and the extended NLM (i.e., MLP+Skip-Gram) separately.⁶ The learned parameters are used to initialize the joint model. All training is conducted by back-propagation using mini-batch SGD (see “Appendix”) to minimize the objective loss (Eq. 1).

Setting β_I & β_T In the unsupervised training, we assign larger β_I to make the training prone to preserve the intramodal semantics of images if the input image feature is of higher quality than the text input feature, and vice versa. For supervised training, since the intramodal semantics are preserved based on reliable labels, we do not distinguish the

⁴ We tried both the Sigmoid function and ReLU activation function for $s()$. ReLU offers better performance.

⁵ Notice that in our model, we fix the word vectors learned by SGM. It can also be fine-tuned by integrating the objective of SGM (Eq. 11) into 15.

⁶ In our experiment, we use the parameters trained by Caff  [18] to initialize the AlexNet to accelerate the training. We use Gensim (<http://radimrehurek.com/gensim/>) to train the skip-gram model with the dimension of word vectors being 100.

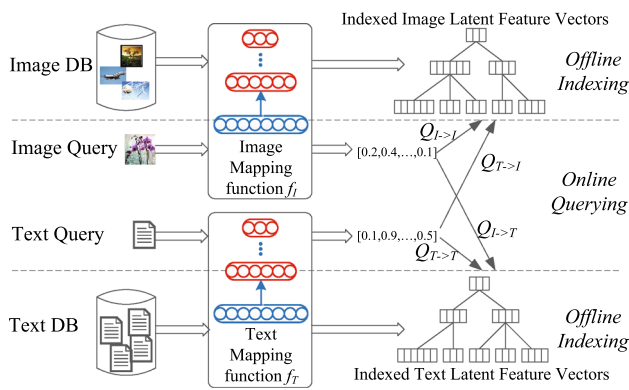


Fig. 8 Illustration of query processing

image modality from the text one in the joint training. Hence, β_I and β_T are set to the same value. In our experiment, to make the three losses within one order of magnitude, we scale the intermodal distance by 0.01.

6 Query processing

After the unsupervised (or supervised) training, each modality has a mapping function. Given a set of heterogeneous data sources, high-dimensional raw features (e.g., bag-of-visual-words or RGB feature for images) are extracted from each source and mapped into a common latent space using the learned mapping functions. In MSAE, we use the image (resp. text) SAE to project image (resp. text) input features into the latent space. In MDNN, then we use the extended DCNN (resp. extended NLM) to map the image (resp. text) input feature into the common latent space.

After the mapping, we create VA-Files [40] over the latent features (one per modality). VA-File is a classic index that can overcome the curse of dimensionality when answering the nearest neighbor queries. It encodes each data point into a bitmap, and the whole bitmap file is loaded into memory for efficient scanning and filtering. Only a small number of real data points will be loaded into memory for verification. Given a query input, we check its media type and map it into the latent space through its modal-specific mapping function. Next, intramodal and intermodal searches are conducted against the corresponding index (i.e., the VA-File) shown in Fig. 8. For example, the task of searching relevant tags of one image, i.e., $Q_{I \rightarrow T}$, is processed by the index for the text latent vectors.

To further improve the search efficiency, we convert the real-valued latent features into binary features, and search based on Hamming distance. The conversion is conducted using existing hash methods that preserve the neighborhood relationship. For example, in our experiment (Sect. 8.2), we use Spectral Hashing [41], which converts real-valued vectors (data points) into binary codes with the objective to

minimize the Hamming distance of data points that are close in the original Euclidean space. Other hashing approaches like [12, 35] are also applicable.

The conversion from real-valued features to binary features trades off effectiveness for efficiency. Since there is information loss when real-valued data are converted to binaries, it affects the retrieval performance. We study the trade-off between efficiency and effectiveness on binary features and real-valued features in the experiment section.

7 Related work

The key problem of multi-modal retrieval is to find an effective mapping mechanism, which maps data from different modalities into a common latent space. An effective mapping mechanism would preserve both intramodal semantics and intermodal semantics well in the latent space and thus generates good retrieval performance.

Linear projection has been studied to solve this problem [21, 36, 44]. The main idea is to find a linear projection matrix for each modality that maps semantic relevant data into similar latent vectors. However, when the distribution of the original data is nonlinear, it would be hard to find a set of effective projection matrices. CVH [21] extends the Spectral Hashing [41] to multi-modal data by finding a linear projection for each modality that minimizes the Euclidean distance of relevant data in the latent space. Similarity matrices for both intermodal data and intramodal data are required to learn a set of good mapping functions. IMH [36] learns the latent features of all training data first before it finds a hash function to fit the input data and output latent features, which could be computationally expensive. LCMH [44] exploits the intramodal correlations by representing data from each modality using its distance to cluster centroids of the training data. Projection matrices are then learned to minimize the distance of relevant data (e.g., image and tags) from different modalities.

Other recent works include CMSSH [4], MLBE [43] and LSCMR [25]. CMSSH uses a boosting method to learn the projection function for each dimension of the latent space. However, it requires prior knowledge such as semantic relevant and irrelevant pairs. MLBE explores correlations of data (both intermodal and intramodal similarity matrices) to learn latent features of training data using a probabilistic graphic model. Given a query, it is converted into the latent space based on its correlation with the training data. Such correlation is decided by labels associated with the query. However, labels of queries are usually not available in practice, which makes it hard to obtain its correlation with the training data. LSCMR [25] learns the mapping functions with the objective to optimize the ranking criteria (e.g., MAP). Ranking examples (a ranking example is a query and its ranking list)

are needed for training. In our algorithm, we use simple relevant pairs (e.g., image and its tags) as training input; thus, no prior knowledge such as irrelevant pairs, similarity matrix, ranking examples and labels of queries is needed.

Multi-modal deep learning [29,37] extends deep learning to multi-modal scenario. [37] combines two Deep Boltzmann Machines (DBM) (one for image and one for text) with a common latent layer to construct a Multi-modal DBM. [29] constructs a bimodal deep auto-encoder with two deep auto-encoders (one for audio and one for video). Both models aim to improve the classification accuracy of objects with features from multiple modalities. They combine different features to learn a (high dimensional) latent feature. In this paper, we aim to represent data with low-dimensional latent features to enable effective and efficient multi-modal retrieval, where both queries and database objects may have features from only one modality. DeVISE [9] from Google shares similar idea with our supervised training algorithm. It embeds image features into text space, which are then used to retrieve similar text features for zero-shot learning. Notice that the text features used in DeVISE to learn the embedding function are generated from high-quality labels. However, in multi-modal retrieval, queries usually do not come with labels and text features are generated from noisy tags. This makes DeVISE less effective in learning robust latent features against noisy input.

8 Experimental study

This section provides an extensive performance study of our solution in comparison with the state-of-the-art methods. We examine both efficiency and effectiveness of our method including training overhead, query processing time and accuracy. Visualization of the training process is also provided to help understand the algorithms. All experiments are conducted on CentOS 6.4 using CUDA 5.5 with NVIDIA GPU (GeForce GTX TITAN). The size of main memory is 64GB and the size of GPU memory is 6GB. The code and hyperparameter settings are available online.⁷ In the rest of this section, we first give our evaluation metrics and then study the performance of unsupervised approach and supervised approach, respectively.

8.1 Evaluation metrics

We evaluate the effectiveness of the mapping mechanism by measuring the accuracy of the multi-modal search, i.e., $\mathbb{Q}_{q \rightarrow t}(q, t \in \{T, I\})$, using the mapped latent features. Without specifications, searches are conducted against real-valued latent features using Euclidean distance. We use the mean

average precision (MAP) [27], one of the standard information retrieval metrics, as the major evaluation metric. Given a set of queries, we calculate the average precision (AP) for each query q as,

$$AP(q) = \frac{\sum_{k=1}^R P(k)\delta(k)}{\sum_{j=1}^R \delta(j)} \quad (17)$$

where R is the size of the test dataset; $\delta(k) = 1$ if the k -th result is relevant, otherwise $\delta(k) = 0$; $P(k)$ is the precision of the result ranked at position k , which is the fraction of true relevant documents in the top k results. By averaging AP for all queries, we get the MAP score. The larger the MAP score, the better the search performance. In addition to MAP, we measure the *precision* and *recall* of search tasks. Given a query, the *ground truth* is defined as: *if a result shares at least one common label (or category) with the query, it is considered as a relevant result; otherwise, it is irrelevant.*

Besides effectiveness, we also evaluate the training overhead in terms of time, cost and memory consumption. In addition, we report query processing time.

8.2 Experimental study of unsupervised approach

First, we describe the datasets used for unsupervised training. Second, an analysis of the training process by visualization is presented. Last, comparison with previous works, including CVH [21], CMSSH [4] and LCMH [44] are provided.⁸

8.2.1 Datasets

Unsupervised training requires relevant image text pairs, which are easy to collect. We use three datasets to evaluate the performance—NUS-WIDE [5], Wiki [30] and Flickr1M [17].

NUS-WIDE The dataset contains 269,648 images from Flickr with each image associated with six tags on average. We refer to the image and its tags as an image–text pair. There are 81 ground truth labels manually annotated for evaluation. Following previous works [24,44], we extract 190,421 image–text pairs annotated with the most frequent 21 labels and split them into three subsets for training, validation and test, respectively. The size of each subset is shown in Table 1. For validation (resp. test), 100 (resp. 1000) queries are randomly selected from the validation (resp. test) dataset. Image and text features are provided in the dataset [5]. For images, SIFT features are extracted and clustered into 500 visual-words. Hence, an image is represented by a 500-dimensional

⁷ <http://www.comp.nus.edu.sg/~wangwei/code>.

⁸ The code and parameter configurations for CVH and CMSSH are available online at <http://www.cse.ust.hk/~dyeyung/code/mlbe.zip>. The code for LCMH is provided by the authors. Parameters are set according to the suggestions provided in the paper.

Table 1 Statistics of datasets for unsupervised training

Dataset	NUS-WIDE	Wiki	Flickr1M
Total size	190,421	2866	1,000,000
Training set	60,000	2000	975,000
Validation set	10,000	366	6000
Test set	120,421	500	6000
Average text length	6	131	5

bag-of-visual-words vector. Its associated tags are represented by a 1000-dimensional tag occurrence vector.

Wiki This dataset contains 2,866 image–text pairs from the Wikipedia’s featured articles. An article in Wikipedia contains multiple sections. The text and its associated image in one section is considered as an image–text pair. Every image–text pair has a label inherited from the article’s category (there are 10 categories in total). We randomly split the dataset into three subsets as shown in Table 1. For validation (resp. test), we randomly select 50 (resp. 100) pairs from the validation (resp. test) set as the query set. Images are represented by 128 dimensional bag-of-visual-words vectors based on SIFT feature. For text, we construct a vocabulary with the most frequent 1,000 words excluding stop words and represent one text section by 1,000 dimensional word count vector like [25]. The average number of words in one section is 131 which is much higher than that in NUS-WIDE. To avoid overflow in Eq. 6 and smooth the text input, we normalize each unit x as $\log(x + 1)$ [32].

Flickr1M This dataset contains 1 million images associated with tags from Flickr, 25,000 of which are annotated with labels (there are 38 labels in total). The image feature is a 3,857 dimensional vector concatenated by SIFT feature, color histogram, etc [37]. Like NUS-WIDE, the text feature is represented by a tag occurrence vector with 2,000 dimensions. All the image–text pairs without annotations are used for training. For validation and test, we randomly select 6,000 pairs with annotations, respectively, among which 1,000 pairs are used as queries.

Before training, we use ZCA whitening [19] to normalize each dimension of image feature to have zero mean and unit variance.

8.2.2 Training visualization

In this section, we visualize the training process of MSAE using the NUS-WIDE dataset as an example to help understand the intuition of the training algorithm and the setting of the weight parameters, i.e., β_I and β_T . Our goal is to learn a set of mapping functions such that the mapped latent features capture both intramodal semantics and intermodal semantics well. Generally, the intermodal semantics is preserved by minimizing the distance of the latent features of

relevant intermodal pairs. The intramodal semantics is preserved by minimizing the reconstruction error of each SAE and through intermodal semantics (see Sect. 4 for details).

First, following the training procedure in Sect. 4, we train a four-layer image SAE with the dimension of each layer as $500 \rightarrow 128 \rightarrow 16 \rightarrow 2$. Similarly, a four-layer text SAE (the structure is $1000 \rightarrow 128 \rightarrow 16 \rightarrow 2$) is trained.⁹ There is no standard guideline for setting the number of latent layers and units in each latent layer for deep learning [1]. In all our experiments, we adopt the widely used pyramid-like structure [6, 15], i.e., decreasing layer size from the bottom (or first hidden) layer to the top layer. In our experiment, we observed that two latent layers perform better than a single latent layer. But there is no significant improvement from two latent layers to three latent layers. Latent features of sampled image–text pairs from the validation set are plotted in Fig. 9a. The pre-training stage initializes SAEs to capture regularities of the original features of each modality in the latent features. On the one hand, the original features may be of low quality to capture intramodal semantics. In such a case, the latent features would also fail to capture the intramodal semantics. We evaluate the quality of the mapped latent features from each SAE by intramodal search on the validation dataset. The MAP of the image intramodal search is about 0.37, while that of the text intramodal search is around 0.51. On the other hand, as the SAEs are trained separately, intermodal semantics are not considered. We randomly pick 25 relevant image–text pairs and connect them with red lines in Fig. 9b. We can see the latent features of most pairs are far away from each other, which indicates that the intermodal semantics are not captured by these latent features. To solve the above problems, we integrate the intermodal loss in the learning objective as Eq. 1. In the following figures, we only plot the distribution of these 25 pairs for ease of illustration.

Second, we adjust the image SAE with the text SAE fixed from epoch 1 to epoch 30. One epoch means one pass of the whole training dataset. Since the MAP of the image intramodal search is worse than that of the text intramodal search, according to the intuition in Sect. 3, we should use a small β_I to decrease the weight of image reconstruction error \mathcal{L}_I in the objective function, i.e., Eq. 1. To verify this, we compare the performance of two choices of β_I , namely $\beta_I = 0$ and $\beta_I = 0.01$. The first two rows in Fig. 10 show the latent features generated by the image SAE after epoch 1 and epoch 30. Comparing image–text pairs in Fig. 10b, d, we can see that with smaller β_I , the image latent features move closer to their relevant text latent features. This is in accordance with Eq. 1, where smaller β_I relaxes the restriction on the image reconstruction error, and in turn increases the

⁹ The last layer with two units is for visualization purpose, such that the latent features could be showed in a 2D space.

Fig. 9 Visualization of latent features after projecting them into 2D space (*Blue points* are image latent features; *White points* are text latent features. Relevant image–text pairs are connected using *red lines*). **a** 300 random image–text pairs. **b** 25 image–text pairs

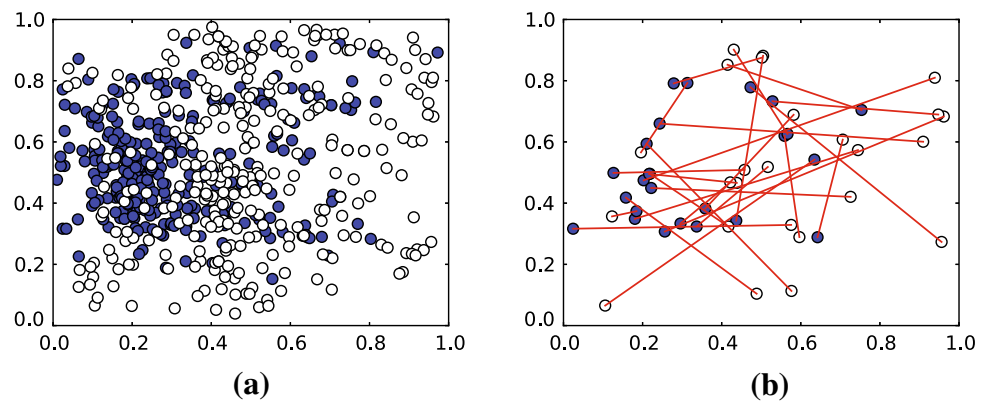


Fig. 10 Adjusting image SAE with different β_I and text SAE fixed (**a–d** show the positions of features of image–text pairs in 2D space). **a** $\beta_I = 0$, epoch 1. **b** $\beta_I = 0$, epoch 30. **c** $\beta_I = 0.01$, epoch 1. **d** $\beta_I = 0.01$, epoch 30. **e** $\beta_I = 0$. **f** $\beta_I = 0.01$

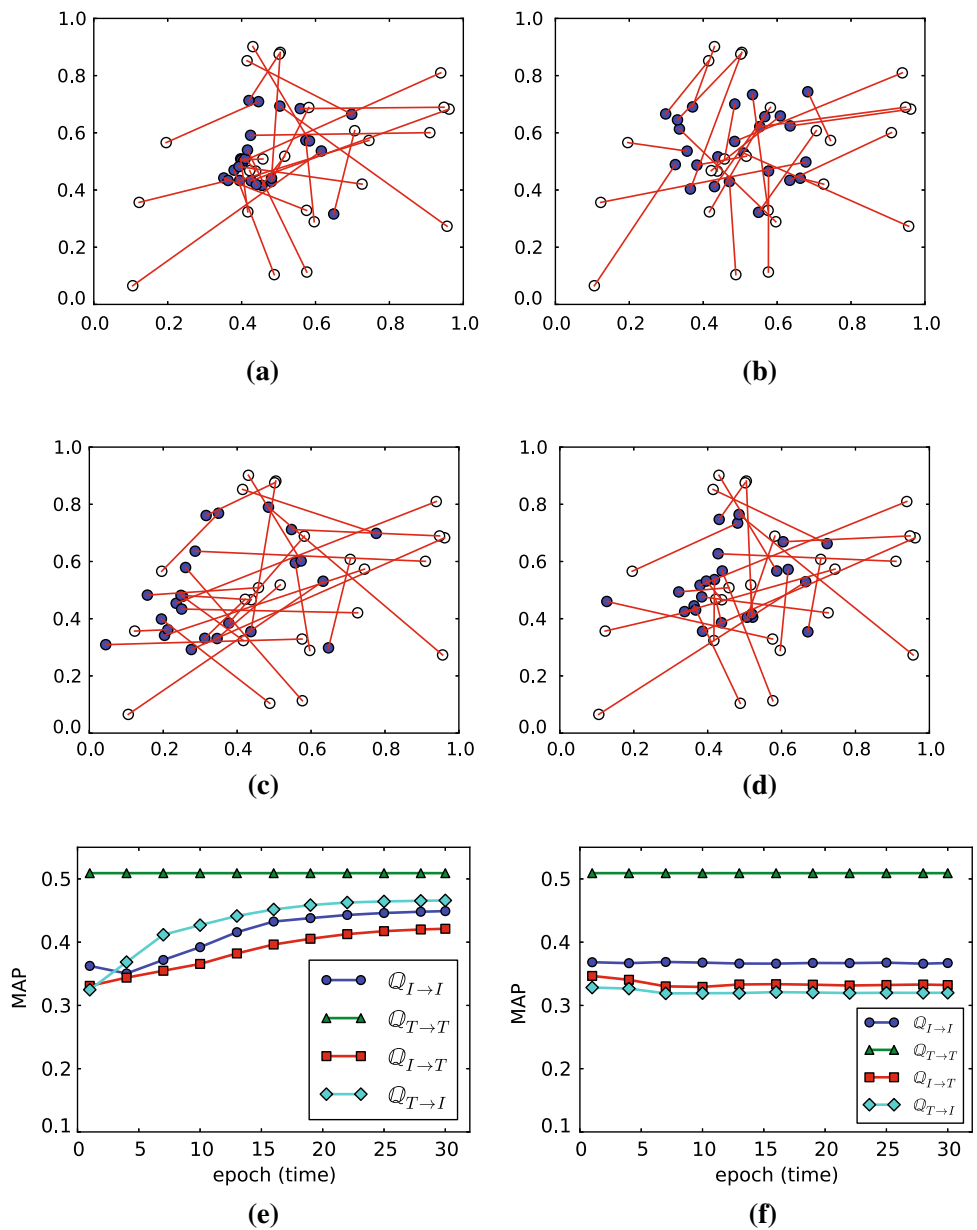
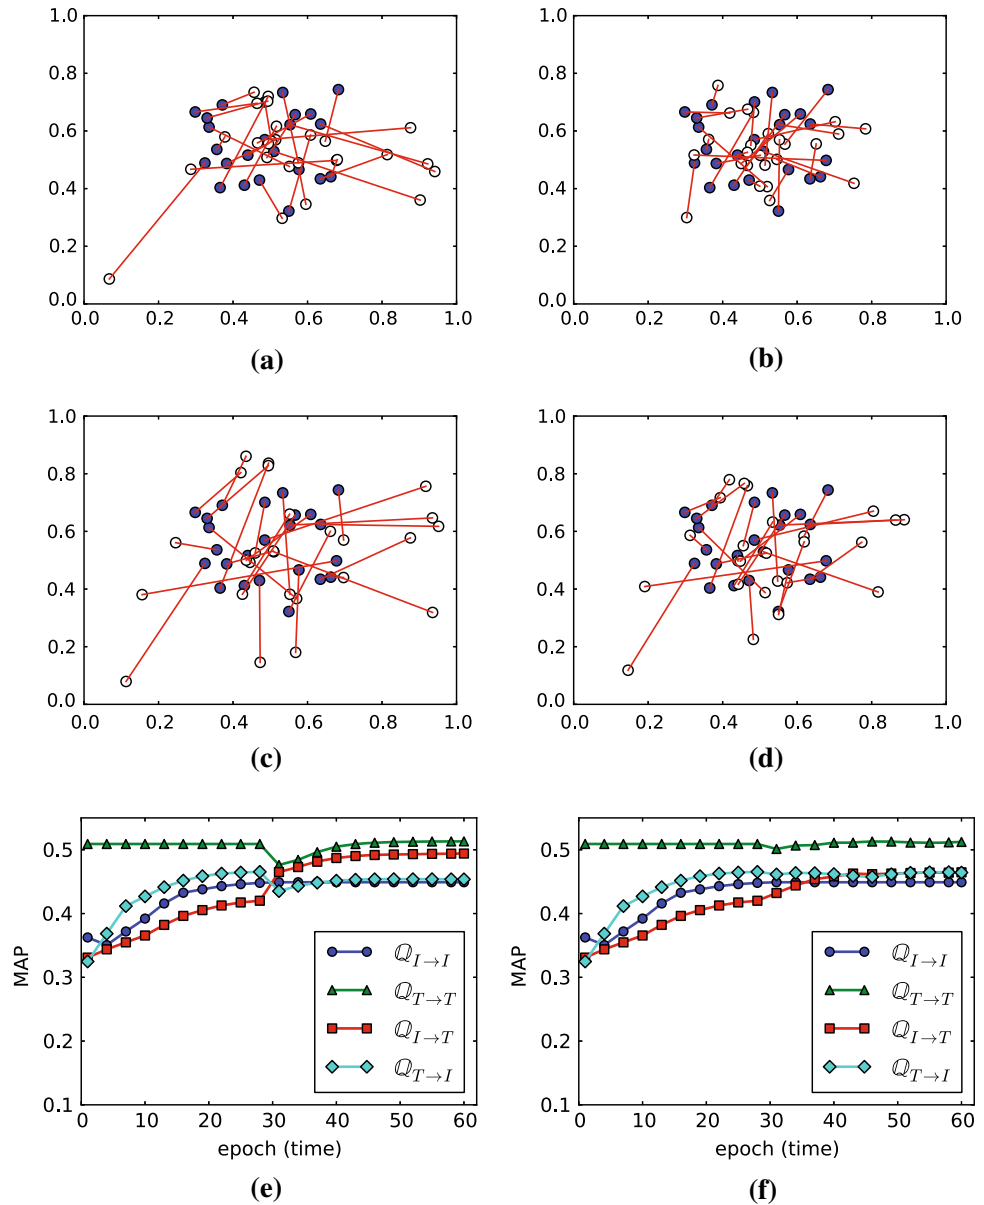


Fig. 11 Adjusting text SAE with different β_T and image SAE fixed (a–d show the positions of features of image–text pairs in 2D space). **a** $\beta_T = 0.01$, epoch 31. **b** $\beta_T = 0.01$, epoch 60. **c** $\beta_T = 0.1$, epoch 31. **d** $\beta_T = 0.1$, epoch 60. **e** $\beta_T = 0.01$. **f** $\beta_T = 0.1$



weight for intermodal distance $\mathcal{L}_{I,T}$. By moving close to relevant text latent features, the image latent features gain more semantics. As shown in Fig. 10e, the MAPs increase as training goes on. MAP of $Q_{T \rightarrow T}$ does not change because the text SAE is fixed. When $\beta_I = 0.01$, the MAPs do not increase in Fig. 10f. This is because image latent features hardly move close to the relevant text latent features as shown in Fig. 10c, d. We can see that the text modality is of better quality for this dataset. Hence it should be assigned a larger weight. However, we cannot set a too large weight for it as explained in the following paragraph.

Third, we adjust the text SAE with the image SAE fixed from epoch 31 to epoch 60. We also compare two choices of β_T , namely 0.01 and 0.1. β_I is set to 0. Figure 11 shows the snapshots of latent features and the MAP curves of each

setting. From Fig. 10b to 11a, which are two consecutive snapshots taken from epoch 30 and 31, respectively, we can see that the text latent features move much closer to the relevant image latent features. It leads to the big changes of MAPs at epoch 31 in Fig. 11e. For example, $Q_{T \rightarrow T}$ substantially drops from 0.5 to 0.46. This is because the sudden moves toward images change the intramodal relationships of text latent features. Another big change happens on $Q_{I \rightarrow T}$, whose MAP increases dramatically. The reason is that when we fix the text features from epoch 1 to 30, an image feature I is pulled to be close to (or nearest neighbor of) its relevant text feature T . However, T may not be the reverse nearest neighbor of I . In epoch 31, we move T toward I such that T is more likely to be the reverse nearest neighbor of I . Hence, the MAP of query $Q_{I \rightarrow T}$ is greatly improved.

On the contrary, $\mathbb{Q}_{T \rightarrow I}$ decreases. From epoch 32 to epoch 60, the text latent features on the one hand move close to relevant image latent features slowly and on the other hand rebuild their intramodal relationships. The latter is achieved by minimizing the reconstruction error \mathcal{L}_T to capture the semantics of the original features. Therefore, both $\mathbb{Q}_{T \rightarrow T}$ and $\mathbb{Q}_{I \rightarrow T}$ grow gradually. Comparing Fig. 11a, c, we can see the distance of relevant latent features in Fig. 11c is larger than that in Fig. 11a. The reason is that when β_T is larger, the objective function in Eq. 1 pays more effort to minimize the reconstruction error \mathcal{L}_T . Consequently, less effort is paid to minimize the intermodal distance $\mathcal{L}_{I,T}$. Hence, relevant intermodal pairs cannot move closer. This effect is reflected as minor changes of MAPs at epoch 31 in Fig. 11f in contrast with that in Fig. 11e. Similarly, small changes happen between Fig. 11c, d, which leads to minor MAP changes from epoch 32 to 60 in Fig. 11f.

8.2.3 Evaluation of model effectiveness on NUS-WIDE dataset

We first examine the mean average precision (MAP) of our method using Euclidean distance against real-valued features. Let L be the dimension of the latent space. Our MSAE is configured with three layers, where the image features are mapped from 500 dimensions to 128 and finally to L . Similarly, the dimension of text features is reduced from 1000 \rightarrow 128 \rightarrow L by the text SAE. β_I and β_T are set to 0 and 0.01, respectively, according to Sect. 8.2.2. We test L with values 16, 24 and 32. The results compared with other methods are reported in Table 2. Our MSAE achieves the best performance for all four search tasks. It demonstrates an average improvement of 17, 27, 21 and 26 % for $\mathbb{Q}_{I \rightarrow I}$, $\mathbb{Q}_{T \rightarrow T}$, $\mathbb{Q}_{I \rightarrow T}$, and $\mathbb{Q}_{T \rightarrow I}$, respectively. CVH and CMSSH prefer smaller L in queries $\mathbb{Q}_{I \rightarrow T}$ and $\mathbb{Q}_{T \rightarrow I}$. The reason is that it needs to train far more parameters in higher dimensions and the learned models will be farther from the optimal solutions. Our method is less sensitive to the value of L . This is probably because with multiple layers, MSAE has stronger representation power and thus is more robust under different L .

Figure 12 shows the precision–recall curves and the recall–candidate ratio curves (used by [43, 44]) which show the change of recall when inspecting more results on the returned rank list. We omit the figures for $\mathbb{Q}_{T \rightarrow T}$ and $\mathbb{Q}_{I \rightarrow I}$ as they show similar trends as $\mathbb{Q}_{T \rightarrow I}$ and $\mathbb{Q}_{I \rightarrow T}$. Our method shows the best accuracy except when recall is 0,¹⁰ whose precision p implies that the nearest neighbor of the query appears in the $\frac{1}{p}$ -th returned result. This indicates that our method performs the best for general top- k similarity retrieval except $k = 1$. For the recall–candidate ratio, the curve of MSAE is always above those of other methods. It shows that we get

¹⁰ Here, recall $r = \frac{1}{\# \text{all relevant results}} \approx 0$.

Table 2 Mean average precision on NUS-WIDE dataset

Task	Algorithm	$\mathbb{Q}_{I \rightarrow I}$				$\mathbb{Q}_{T \rightarrow T}$				$\mathbb{Q}_{I \rightarrow T}$				$\mathbb{Q}_{T \rightarrow I}$			
		LCMH	CMSSH	CVH	MSAE	LCMH	CMSSH	CVH	MSAE	LCMH	CMSSH	CVH	MSAE	LCMH	CMSSH	CVH	MSAE
Dimension of latent space L	16	0.353	0.355	0.365	0.417	0.373	0.400	0.374	0.498	0.328	0.391	0.359	0.447	0.331	0.337	0.368	0.432
	24	0.343	0.356	0.358	0.412	0.373	0.402	0.364	0.480	0.333	0.388	0.351	0.444	0.323	0.336	0.360	0.427
	32	0.343	0.357	0.354	0.413	0.374	0.403	0.357	0.470	0.333	0.382	0.345	0.402	0.324	0.335	0.355	0.435

The bold numbers are the best scores

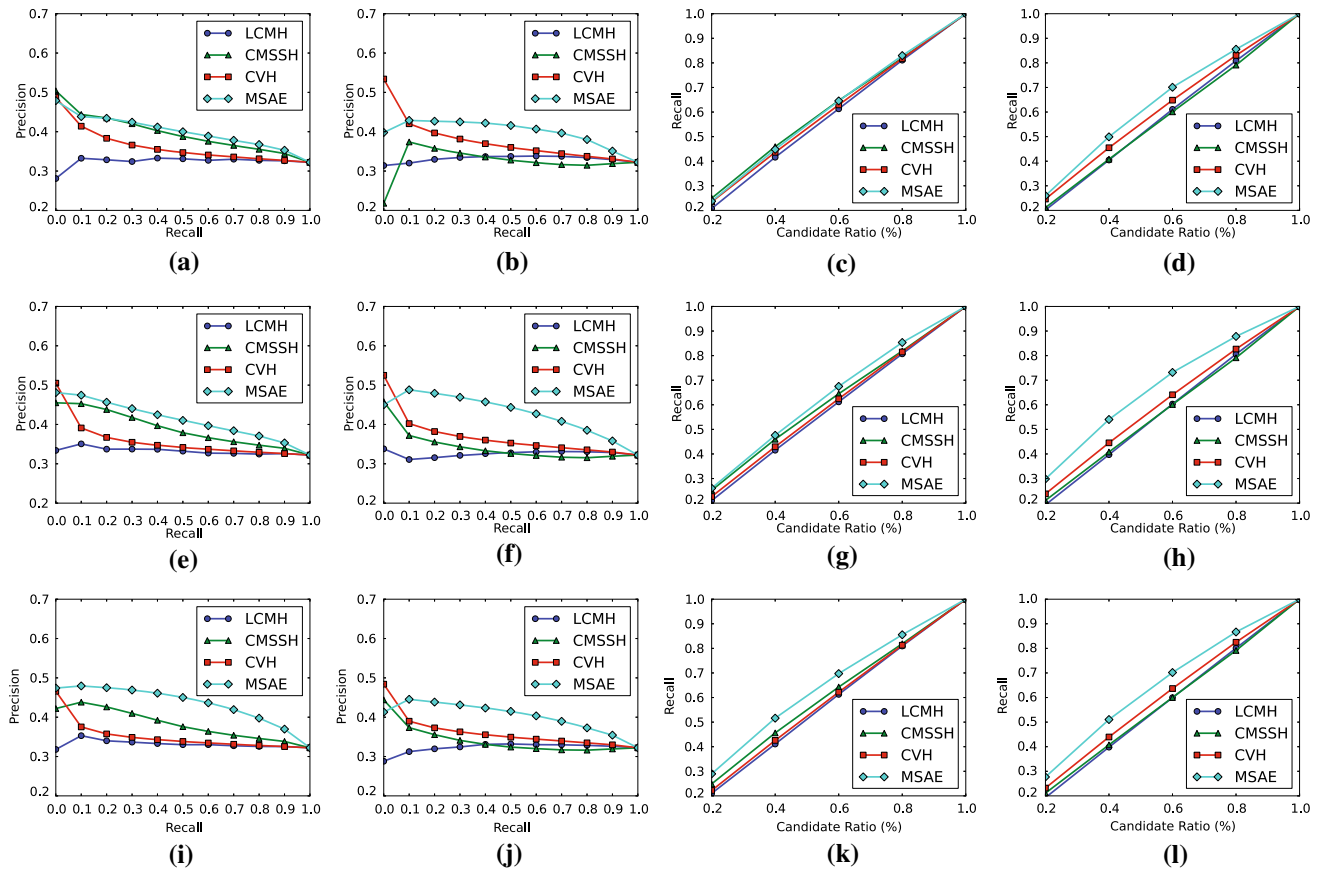


Fig. 12 Precision–recall and recall–candidate ratio on NUS-WIDE dataset. **a** $\mathcal{Q}_{I \rightarrow T}$, $L = 16$. **b** $\mathcal{Q}_{T \rightarrow I}$, $L = 16$. **c** $\mathcal{Q}_{I \rightarrow T}$, $L = 16$. **d** $\mathcal{Q}_{T \rightarrow I}$, $L = 16$. **e** $\mathcal{Q}_{I \rightarrow T}$, $L = 24$. **f** $\mathcal{Q}_{T \rightarrow I}$, $L = 24$. **g** $\mathcal{Q}_{I \rightarrow T}$, $L = 24$. **h** $\mathcal{Q}_{T \rightarrow I}$, $L = 24$. **i** $\mathcal{Q}_{I \rightarrow T}$, $L = 32$. **j** $\mathcal{Q}_{T \rightarrow I}$, $L = 32$. **k** $\mathcal{Q}_{I \rightarrow T}$, $L = 32$. **l** $\mathcal{Q}_{T \rightarrow I}$, $L = 32$

better recall when inspecting the same number of objects. In other words, our method ranks more relevant objects at higher (front) positions. Therefore, MSAE performs better than other methods.

Besides real-valued features, we also conduct experiments against binary latent features for which Hamming distance is used as the distance function. In our implementation, we choose Spectral Hashing [41] to convert real-valued latent feature vectors into binary codes. Other comparison algorithms use their own conversion mechanisms. The MAP scores are reported in Table 3. We can see that 1) MSAE still performs better than other methods. 2) The MAP scores using Hamming distance are not as good as those of Euclidean distance. This is due to the possible information loss by converting real-valued features into binary features.

8.2.4 Evaluation of model effectiveness on Wiki dataset

We conduct similar evaluations on **Wiki** dataset as on NUS-WIDE. For MSAE with latent feature of dimension L , the structure of its image SAE is $128 \rightarrow 128 \rightarrow L$, and the

structure of its text SAE is $1000 \rightarrow 128 \rightarrow L$. Similar to the settings on NUS-WIDE, β_I is set to 0 and β_T is set to 0.01.

The performance is reported in Table 4. MAPs on **Wiki** dataset are much smaller than those on NUS-WIDE except for $\mathcal{Q}_{T \rightarrow T}$. This is because the images of **Wiki** are of much lower quality. It contains only 2000 images that are highly diversified, making it difficult to capture the semantic relationships among images, and between images and text. Query task $\mathcal{Q}_{T \rightarrow T}$ is not affected as Wikipedia’s featured articles are well edited and rich in text information. In general, our method achieves an average improvement of 8.1, 30.4, 32.8 and 26.8 % for $\mathcal{Q}_{I \rightarrow I}$, $\mathcal{Q}_{T \rightarrow T}$, $\mathcal{Q}_{I \rightarrow T}$ and $\mathcal{Q}_{T \rightarrow I}$, respectively. We do not plot the precision–recall curves and recall–candidate ratio curves as they show similar trends to those of NUS-WIDE.

8.2.5 Evaluation of model effectiveness on Flickr1M dataset

We configure a four-layer image SAE as $3857 \rightarrow 1000 \rightarrow 128 \rightarrow L$ and a four-layer text SAE as $2000 \rightarrow 1000 \rightarrow 128 \rightarrow L$ for this dataset. Different from the other two datasets, the original image feature of Flickr1M is of higher

Table 3 Mean average precision on NUS-WIDE dataset (using binary latent features)

Task	$\mathbb{Q}_{I \rightarrow I}$					$\mathbb{Q}_{I \rightarrow T}$					$\mathbb{Q}_{T \rightarrow I}$				
	LCMH	CMSSH	CVH	MSAE		LCMH	CMSSH	CVH	MSAE		LCMH	CMSSH	CVH	MSAE	
Dimension of latent space L	16	0.353	0.357	0.352	0.376	0.387	0.391	0.379	0.397	0.328	0.339	0.359	0.364	0.325	0.346
	24	0.347	0.358	0.346	0.368	0.392	0.396	0.372	0.412	0.333	0.346	0.353	0.371	0.324	0.352
	32	0.345	0.358	0.343	0.359	0.395	0.397	0.365	0.434	0.320	0.340	0.348	0.373	0.318	0.347

The bold numbers are the best scores

Table 4 Mean average precision on Wiki dataset

Task	$\mathbb{Q}_{I \rightarrow I}$					$\mathbb{Q}_{I \rightarrow T}$					$\mathbb{Q}_{T \rightarrow I}$				
	LCMH	CMSSH	CVH	MSAE		LCMH	CMSSH	CVH	MSAE		LCMH	CMSSH	CVH	MSAE	
Dimension of latent space L	16	0.146	0.148	0.147	0.162	0.359	0.318	0.153	0.462	0.133	0.138	0.126	0.182	0.117	0.140
	24	0.149	0.151	0.150	0.161	0.345	0.320	0.151	0.437	0.129	0.135	0.123	0.176	0.124	0.138
	32	0.147	0.149	0.148	0.162	0.333	0.312	0.152	0.453	0.137	0.133	0.128	0.187	0.119	0.137

The bold numbers are the best scores

quality as it consists of both local and global features. For intramodal search, the image latent feature performs equally well as the text latent feature. Therefore, we set both β_I and β_T to 0.01.

We compare the MAP of MSAE and CVH in Table 5. MSAE outperforms CVH in most of the search tasks. The results of LCMH and CMSSH cannot be reported as both methods run out of memory in the training stage.

8.2.6 Evaluation of training cost

We use the largest dataset Flickr1M to evaluate the training cost of time and memory consumption. The results are reported in Fig. 13. The training cost of LCMH and CMSSH is not reported because they run out of memory on this dataset. We can see that the training time of MSAE and CVH increases linearly with respect to the size of the training dataset. Due to the stacked structure and multiple iterations of passing the dataset, MSAE is not as efficient as CVH. Roughly, the overhead is about the number of training iterations times the height of MSAE. Possible solutions for accelerating the MSAE training include adopting distributed deep learning [7]. We leave this as our future work.

Figure 13b shows the memory usage of the training process. Given a training dataset, MSAE splits them into mini-batches and conducts the training batch by batch. It stores the model parameters and one mini-batch in memory, both of which are independent of the training dataset size. Hence, the memory usage stays constant when the size of the training dataset increases. The actual minimum memory usage for MSAE is smaller than 10GB. In our experiments, we allocate more space to load multiple mini-batches into memory to save disk reading cost. CVH has to load all training data into memory for matrix operations. Therefore, its memory usage increases with respect to the size of the training dataset.

8.2.7 Evaluation of query processing efficiency

We compare the efficiency of query processing using binary latent features and real-valued latent features. Notice that all methods (i.e., MSAE, CVH, CMSSH and LCMH) perform similarly in query processing after mapping the original data into latent features of same dimension. Data from the Flickr1M training dataset are mapped into a 32 dimensional latent space to form a large dataset for searching. To speed up the query processing of real-valued latent features, we create an index (i.e., VA-File [40]) for each modality. For binary latent features, we do not create any indexes, as linear scan offers decent performance as shown in Figure 14. It shows the time of searching 50 nearest neighbors (averaged over 100 random queries) against datasets represented using binary latent features (based on Hamming distance) and real-valued

features (based on Euclidean distance), respectively. We can see that the querying time increases linearly with respect to the dataset size for both binary and real-valued latent features. But the searching against binary latent features is 10× faster than that against real-valued latent features. This is because the computation of Hamming distance is more efficient than that of Euclidean distance.

By taking into account the results from effectiveness evaluations, we can see that there is a trade-off between efficiency and effectiveness in feature representation. The binary encoding greatly improves the efficiency in the expense of accuracy degradation (Table 3).

8.3 Experimental study of supervised approach

8.3.1 Datasets

Supervised training requires input image–text pairs to be associated with additional semantic labels. Since Flickr1M does not have labels and Wiki dataset has too few labels that are not discriminative enough, we use NUS-WIDE dataset to evaluate the performance of supervised training. We extract 203,400 labeled pairs, among which 150,000 are used for training. The remaining pairs are evenly partitioned into two sets for validation and testing. From both sets, we randomly select 2000 pairs as queries. This labeled dataset is named NUS-WIDE-a.

We further extract another dataset from NUS-WIDE-a by filtering those pairs with more than one label. This dataset is named NUS-WIDE-b and is used to compare with DeViSE [9], which is designed for training against images annotated with single label. In total, we obtain 76,000 pairs. Among them, we randomly select 60,000 pairs for training, and the rest are evenly partitioned for validation and testing. Thousand queries are randomly selected from the two datasets, respectively. The two datasets are summarized in Table 6.


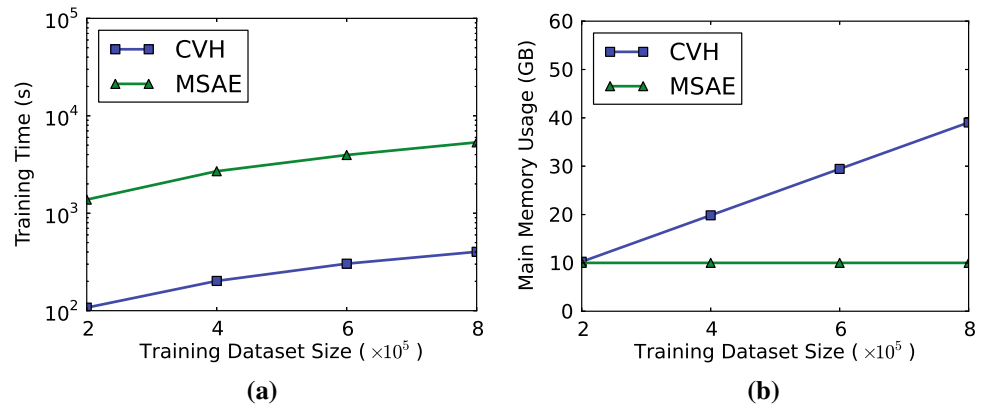
8.3.2 Visualization of training process

NUS-WIDE-a In Fig. 15a, we plot the total training loss \mathcal{L} and its components (\mathcal{L}_I , \mathcal{L}_T and $\mathcal{L}_{I,T}$) in the first 50,000 iterations (one iteration for one mini-batch) against the NUS-WIDE-a dataset. We can see that training converges rather quickly. The training loss drops dramatically at the very beginning and then decreases slowly. This is because initially the learning rate is large and the parameters approach quickly toward the optimal values. Another observation is that the intramodal loss \mathcal{L}_I for the image modality is smaller than \mathcal{L}_T for the text modality. This is because some tags may be noisy or not very relevant to the associated labels that represent the main visual content in the images. It is difficult to learn a set of parameters to map noisy tags into the

Table 5 Mean average precision on Flickr1M dataset

Task		$Q_{I \rightarrow I}$		$Q_{T \rightarrow T}$		$Q_{I \rightarrow T}$		$Q_{T \rightarrow I}$	
Algorithm		CVH	MSAE	CVH	MSAE	CVH	MSAE	CVH	MSAE
L	16	0.622	0.621	0.610	0.624	0.610	0.632	0.616	0.608
	24	0.616	0.619	0.604	0.629	0.605	0.628	0.612	0.612
	32	0.603	0.622	0.587	0.630	0.588	0.632	0.598	0.614

The bold numbers are the best scores

Fig. 13 Training cost comparison on Flickr1M dataset**Fig. 14** Querying time comparison using real-valued and binary latent features

latent space and well predict the ground truth labels. The intermodel training loss is calculated at a different scale and is normalized to be within one order of magnitude as \mathcal{L}_I and \mathcal{L}_T .

The MAPs for all types of searches using supervised training model are shown in Fig. 15b. As can be seen, the MAPs first gradually increase and then become stable in the last few iterations. It is worth noting that the MAPs are much higher than the results of unsupervised training (MSAE) in Fig. 11. There are two reasons for the superiority. First, the supervised training algorithm (MDNN) exploits DCNN and NLM to learn better visual and text features, respectively. Second, labels bring in more semantics and make latent features more robust to noises in input data (e.g., visual irrelevant tags).

Besides MAP, we also evaluate MDNN for multi-label prediction based on precision and recall. For each image (or text), we look at its labels with the largest k probabilities

Table 6 Statistics of datasets for supervised training

Dataset	NUS-WIDE-a	NUS-WIDE-b
Total size	203,400	76,000
Training set	150,000	60,000
Validation set	26,700	80,000
Test set	26,700	80,000

based on Eq. 9 (or 16). For the i -th image (or text), let N_i denote the number of labels out of k that belong to its ground truth label set, and T_i the size of its ground truth label set. The precision and recall are defined according to [11] as follows:

$$\text{precision} = \frac{\sum_{i=1}^n N_i}{k * n}, \quad \text{recall} = \frac{\sum_{i=1}^n N_i}{\sum_{i=1}^n T_i} \quad (18)$$

where n is the test set size. The results are shown in Fig. 15c ($k = 3$). The performance decreases at the early stage and then goes up. This is because at the early stage, in order to minimize the intermodal loss, the training may disturb the pre-trained parameters fiercely that affects the intramodal search performance. Once the intermodal loss is reduced to a certain level, it starts to adjust the parameters to minimize both intermodal loss and intramodal loss. Hence, the classification performance starts to increase. We can also see that the performance of latent text features is not as good as that of latent image features due to the noises in tags. We use the same experiment setting as that in [11], the (over all) precision and recall are 7 and 7.5% higher than that in [11], respectively.




Fig. 15 Visualization of training on NUS-WIDE-a. **a** Training loss. **b** MAP on validation dataset. **c** Precision–recall on validation dataset





Fig. 16 Visualization of training on NUS-WIDE-b. **a** Training loss. **b** MAP on validation dataset. **c** Precision–recall on validation dataset

Fig. 17 Visualization of latent features learned by MDNN for the test dataset of NUS-WIDE-a (features represented by the same shapes and colors are annotated with the same label). **a** Image latent feature. **b** Text latent feature



NUS-WIDE-b Figure 16 shows the training results against the NUS-WIDE-b dataset. The results demonstrate similar patterns to those in Fig. 15. However, MAPs become lower, possibly due to smaller training dataset size and fewer number of associated labels. In Fig. 16c, the precision and recall for classification using image (or text) latent features are the same. This is because each image–text pair has only one label and $T_i = 1$. When we set $k = 1$, the denominator $k * n$ in precision is equal to $\sum_{i=1}^n T_i$ in recall.

2D Visualization To demonstrate that the learned mapping functions can generate semantic discriminative latent features, we extract top-8 most popular labels, and for each

label, we randomly sample 300 image–text pairs from the test dataset of NUS-WIDE-b. Their latent features are projected into a two-dimensional space by t-SNE [26]. Figure 17a shows the two-dimensional image latent features where one point represents one image feature and Fig. 17b shows the two-dimensional text features. Labels are distinguished using different shapes. We can see that the features are well clustered according to their labels. Further, the image features and text features semantically relevant to the same labels are projected to similar positions in the 2D space. For example, in both figures, the red circles are at the left side, and the blue right triangles are in the top area. The two figures together

confirm that our supervised training is very effective in capturing semantic information for multi-modal data.

8.3.3 Evaluation of model effectiveness on NUS-WIDE dataset

In our final experiment, we compare MDNN with DeViSE [9] in terms of effectiveness of multi-modal retrieval. DeViSE maps image features into text feature space. The learning objective is to minimize the rank hinge loss based on the latent features of an image and its labels. We implement this algorithm and extend it to handle multiple labels by averaging their word vector features. We denote this algorithm as DeViSE-L. Besides, we also implement a variant of DeViSE denoted as DeViSE-T, whose learning objective is to minimize the rank hinge loss based on the latent features of an image and its tag(s). Similarly, if there are multiple tags, we average their word vectors. The results are shown in Table 7. The retrieval is conducted using real-valued latent feature and cosine similarity as the distance function. We can see that MDNN performs much better than both DeViSE-L and DeViSE-T for all four types of searches on both NUS-WIDE-a and NUS-WIDE-b. The main reason is that the image tags are not all visually relevant, which makes it hard for the text (tag) feature to capture the visual semantics in DeViSE. MDNN exploits the label information in the training, which helps to train a model that can generate more robust feature against noisy input tags. Hence, the performance of MDNN is better.

8.3.4 Evaluation of training cost

We report the training cost in terms of training time (Fig. 18a) and memory consumption (Fig. 18b) on NUS-WIDE-a dataset. Training time includes the pre-training for each single modality and the joint multi-modal training. MDNN and DeViSE-L take longer time to train than MSAE, because the convolution operations in them are time-consuming. Further, MDNN involves pre-training stages for the image modality and text modality and thus incurs longer training time than DeViSE-L. The memory footprint of MDNN is similar to that of DeViSE-L, as the two methods both rely on DCNN, which incurs most of the memory consumption. DeViSE-L uses features of higher dimension (100 dimension) than MDNN (81 dimension), which leads to about 100 megabytes difference as shown in Fig. 18b.

8.3.5 Comparison with unsupervised approach


By comparing Tables 7 and 2, we can see that the supervised approach—MDNN, performs better than the unsupervised approach—MSAE. This is not surprising because MDNN consumes more information than MSAE. Although the two

Table 7 Mean average precision using real-valued latent feature

Task	Algorithm	$Q_{I \rightarrow I}$		$Q_{T \rightarrow T}$		$Q_{I \rightarrow T}$		$Q_{T \rightarrow I}$	
		MDNN	DeViSE-L	DeViSE-T	MDNN	DeViSE-L	DeViSE-T	MDNN	DeViSE-L
Dataset	NUS-WIDE-a	0.669	0.5619	0.5399	0.541	0.468	0.464	0.612	0.502
	NUS-WIDE-b	0.556	0.432	0.419	0.466	0.367	0.385	0.495	0.222
						0.270	0.399		0.406

The bold numbers are the best scores

Fig. 18 Training cost comparison on NUS-WIDE-a dataset. **a** Training time. **b** Memory consumption



methods share the same general training objective, the exploitation of label semantics helps MDNN learn better features in capturing the semantic relevance of the data from different modalities. For memory consumption, MDNN and MSAE perform similarly (Fig. 18b).

9 Conclusion

In this paper, we have proposed a general framework (objective) for learning mapping functions for effective multi-modal retrieval. Both intramodal and intermodal semantic relationships of data from heterogeneous sources are captured in the general learning objective function. Given this general objective, we have implemented one unsupervised training algorithm and one supervised training algorithm separately to learn the mapping functions based on deep learning techniques. The unsupervised algorithm uses stacked auto-encoders as the mapping functions for the image modality and the text modality. It only requires simple image-text pairs for training. The supervised algorithm uses an extend DCNN as the mapping function for images and an extend NLM as the mapping function for text data. Label information is integrated in the training to learn robust mapping functions against noisy input data. The results of experiment confirm the improvements of our method over previous works in search accuracy. Based on the processing strategies outlined in this paper, we have built a distributed training platform (called SINGA) to enable efficient deep learning training that supports training large-scale deep learning models. We shall report the system architecture and its performance in a future work.

Acknowledgments This work is supported by A*STAR Project 1321202073. Xiaoyan Yang is supported by Human-Centered Cyber-physical Systems (HCCS) programme by A*STAR in Singapore.

Appendix

Algorithm 1 MiniBatchSGD(D, b, f, α)

Input: D , training dataset
Input: b , batchsize
Input: f , the initial mapping function
Input: α , learning rate//set manually
Output: f , updated mapping function.

1. $\theta \leftarrow f$ //initial parameters of f
2. **repeat**
3. **for** $i = 0$ to $|D|/b$ **do**
4. $B = D[i * b : i * b + b]$
5. $\nabla\theta = \text{average}(\text{BackPropagation}(x, f) \text{ for } x \in B)$
6. $\theta = \theta + \alpha * \nabla\theta$
7. **until** Converge
8. **return** f

Algorithm 2 BackPropagation(x_0, f)

Input: x_0 , input feature vector
Input: f , the mapping function with parameter set θ
Output: θ , gradients of parameters of f

1. $h \leftarrow$ height of layers in f
2. $\{x_i\}_{i=1}^h \leftarrow f(x_0|\theta)$ // forward through all layers
3. $\delta_h = \frac{\partial \mathcal{L}}{\partial x_h}$
4. **for** $i = h$ to 1 **do**
5. $\nabla\theta_i = \delta_i * \frac{\partial x_i}{\partial \theta_i} // \frac{\partial \mathcal{L}}{\partial \theta_i}$
6. $\delta_{i-1} = \delta_i * \frac{\partial x_i}{\partial x_{i-1}}$
7. **return** $\{\nabla\theta_i\}_{i=1}^h$

In this section, we present the mini-batch stochastic gradient descent (mini-batch SGD) algorithm and the back-propagation (BP) algorithm [22], which are used throughout this paper to train MSAE and MDNN.

Mini-batch SGD minimizes the objective loss (e.g., $\mathcal{L}, \mathcal{L}_I, \mathcal{L}_T$) by updating the parameters involved in the mapping function(s) based on the gradients of the objective w.r.t the parameters. Specifically, it iterates the whole dataset to

extract mini-batches (Line 4). For each mini-batch, it averages the gradients computed from BP (Line 5) and updates the parameters (Line 6).

BP calculates the gradients of the objective loss (e.g., \mathcal{L} , \mathcal{L}_I , \mathcal{L}_T) w.r.t. the parameters involved in the mapping function (e.g., f_I , f_T) using a chain rule (Eqs. 19, 20). It forwards the input feature vector through all layers of the mapping function (Line 2). Then it backwards the gradients according to the chain rule (Line 4–6). θ_i denotes parameters involved in the i -th layer. Gradients are returned at Line 7.

$$\frac{\partial \mathcal{L}}{\partial \theta_i} = \frac{\partial \mathcal{L}}{\partial x_i} * \frac{\partial x_i}{\partial \theta_i} \quad (19)$$

$$\frac{\partial \mathcal{L}}{\partial x_{i-1}} = \frac{\partial \mathcal{L}}{\partial x_i} * \frac{\partial x_i}{\partial x_{i-1}} \quad (20)$$

References

- Bengio, Y.: Practical recommendations for gradient-based training of deep architectures. CoRR [arXiv:1206.5533](#) (2012)
- Bengio, Y., Ducharme, R., Vincent, P., Janvin, C.: A neural probabilistic language model. *J. Mach. Learn. Res.* **3**, 1137–1155 (2003)
- Bengio, Y., Courville, A.C., Vincent, P.: Representation learning: a review and new perspectives. *IEEE Trans. Pattern Anal. Mach. Intell.* **35**(8), 1798–1828 (2013)
- Bronstein, M.M., Bronstein, A.M., Michel, F., Paragios, N.: Data fusion through cross-modality metric learning using similarity-sensitive hashing. In: *The Twenty-Third IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2010, San Francisco, CA, USA*, pp. 3594–3601. IEEE Computer Society (2010)
- Chua, T.S., Tang, J., Hong, R., Li, H., Luo, Z., Zheng, Y.T.: Nus-wide: a real-world web image database from national university of singapore. In: *Proceedings of ACM Conference on Image and Video Retrieval (CIVR'09)*, Santorini, Greece (2009)
- Ciresan, D.C., Meier, U., Gambardella, L.M., Schmidhuber, J.: Deep Big Multilayer Perceptrons for Digit Recognition, vol. 7700. Springer, Berlin (2012)
- Dean, J., Corrado, G., Monga, R., Chen, K., Devin, M., Le, Q.V., Mao, M.Z., Ranzato, M., Senior, A.W., Tucker, P.A., Yang, K., Ng, A.Y.: Large scale distributed deep networks. In: Bartlett, P.L., Pereira, F.C.N., Burges, C.J.C., Weinberger, K.Q. (eds.) *Advances in Neural Information Processing Systems 25: 26th Annual Conference on Neural Information Processing Systems 2012, Lake Tahoe, Nevada, United States*, pp. 1232–1240 (2012)
- Donahue, J., Jia, Y., Vinyals, O., Hoffman, J., Zhang, N., Tzeng, E., Darrell, T.: Decaf: a deep convolutional activation feature for generic visual recognition. [arXivpreprint arXiv:1310.1531](#) (2013)
- Frome, A., Corrado, G.S., Shlens, J., Bengio, S., Dean, J., Ranzato, M., Mikolov, T.: Devise: a deep visual-semantic embedding model. In: Burges, C.J.C., Bottou, L., Ghahramani, Z., Weinberger, K.Q. (eds.) *Advances in Neural Information Processing Systems 26: 27th Annual Conference on Neural Information Processing Systems 2013, Lake Tahoe, Nevada, United States*, pp. 2121–2129 (2013)
- Girshick, R.B., Donahue, J., Darrell, T., Malik, J.: Rich feature hierarchies for accurate object detection and semantic segmentation. In: *CVPR 2014, Columbus, OH, USA*, pp. 580–587 (2014)
- Gong, Y., Jia, Y., Leung, T., Toshev, A., Ioffe, S.: Deep convolutional ranking for multilabel image annotation. CoRR [arXiv:1312.4894](#) (2013a)
- Gong, Y., Lazebnik, S., Gordo, A., Perronnin, F.: Iterative quantization: a procrustean approach to learning binary codes for large-scale image retrieval. *IEEE Trans. Pattern Anal. Mach. Intell.* **35**(12), 2916–2929 (2013b)
- Goroshin, R., LeCun, Y.: Saturating auto-encoder. CoRR [arXiv:1301.3577](#) (2013)
- Hinton, G.: A Practical Guide to Training Restricted Boltzmann Machines. In: Montavon, G., Müller, K.-R. (eds.) *Neural Networks: Tricks of the Trade-Second Edition, Lecture Notes in Computer Science*, vol. 7700, pp. 599–619. Springer (2012)
- Hinton, G., Salakhutdinov, R.: Reducing the dimensionality of data with neural networks. *Science* **313**(5786), 504–507 (2006)
- Hjaltason, G.R., Samet, H.: Index-driven similarity search in metric spaces. *ACM Trans. Database Syst.* **28**(4), 517–580 (2003)
- Huiskes, M.J., Lew, M.S.: The mir flickr retrieval evaluation. In: *Proceedings of the 1st ACM International Conference on Multimedia Information Retrieval, MIR '08, Vancouver, British Columbia, Canada*, pp. 39–43. ACM, New York, USA (2008)
- Jia, Y., Shelhamer, E., Donahue, J., Karayev, S., Long, J., Girshick, R.B., Guadarrama, S., Darrell, T.: Caffe: Convolutional architecture for fast feature embedding. In: Hua, K.A., Rui, Y., Steinmetz, R., Hanjalic, A., Natsev, A., Zhu, W. (eds.) *Proceedings of the ACM International Conference on Multimedia, MM '14, Orlando, FL, USA*, pp. 675–678. ACM (2014)
- Krizhevsky, A.: Learning Multiple Layers of Features from Tiny Images. Tech. rep (2009)
- Krizhevsky, A., Sutskever, I., Hinton, G.E.: Imagenet classification with deep convolutional neural networks. *Adv. Neural Inf. Process. Syst.* **25**, 1106–1114 (2012)
- Kumar, S., Udupa, R.: Learning hash functions for cross-viewsimilarity search. In: Walsh, T. (ed.) *Proceedings of the 22nd International Joint Conference on Artificial Intelligence, Barcelona, Catalonia*, pp. 1360–1365. IJCAI/AAAI (2011)
- LeCun, Y., Bottou, L., Orr, G., Müller, K.: Efficient backProp. In: Orr, G., Müller, K.R. (eds.) *Neural Networks: Tricks of the Trade, Lecture Notes in Computer Science*, chap 2, vol. 1524, pp. 9–50. Springer, Berlin (1998)
- Liu, D., Hua, X., Yang, L., Wang, M., Zhang, H.: Tag ranking. In: *Proceedings of the 18th International Conference on World Wide Web, WWW 2009, Madrid, Spain*, pp. 351–360, (2009). doi:[10.1145/1526709.1526757](#)
- Liu, W., Wang, J., Kumar, S., Chang, S.F.: Hashing with graphs. In: Getoor, L., Scheffer, T. (eds.) *Proceedings of the 28th International Conference on Machine Learning, ICML 2011, Bellevue, Washington*, pp. 1–8. Omnipress (2011)
- Lu, X., Wu, F., Tang, S., Zhang, Z., He, X., Zhuang, Y.: A low rank structural large margin method for cross-modal ranking. In: *SIGIR*, pp. 433–442 (2013)
- van der Maaten, L.: Accelerating t-SNE using tree-based algorithms. *J. Mach. Learn. Res.* **15**, 3221–3245 (2014)
- Manning, C.D., Raghavan, P., Schütze, H.: *Introduction to Information Retrieval*. Cambridge University Press, Cambridge (2008)
- Ngiam, J., Khosla, A., Kim, M., Nam, J., Lee, H., Ng, A.Y.: Multimodal deep learning. In: Getoor, L., Scheffer, T. (eds.) *Proceedings of the 28th International Conference on Machine Learning, ICML 2011, Bellevue, Washington*, pp. 689–696. Omnipress (2011)
- Ngiam, J., Khosla, A., Kim, M., Nam, J., Lee, H., Ng, A.Y.: Multimodal deep learning. In: Getoor, L., Scheffer, T. (eds.) *Proceedings of the 28th International Conference on Machine Learning, ICML 2011, Bellevue, Washington*, pp. 689–696. Omnipress (2011)
- Rasiwasia, N., Pereira, J.C., Coviello, E., Doyle, G., Lanckriet, G.R.G., Levy, R., Vasconcelos, N.: A new approach to cross-modal multimedia retrieval. In: *ACM Multimedia*, pp. 251–260 (2010)
- Rifai, S., Vincent, P., Muller, X., Glorot, X., Bengio, Y.: Contractive auto-encoders: explicit invariance during feature extraction. In: *ICML*, pp. 833–840 (2011)

32. Salakhutdinov, R., Hinton, G.E.: Semantic hashing. *Int. J. Approx. Reason.* **50**(7), 969–978 (2009)
33. Socher, R., Manning, C.D.: Deep learning for NLP (without magic). In: *Human Language Technologies: Conference of the North American Chapter of the Association of Computational Linguistics, Proceedings*, pp. 1–3. Westin Peachtree Plaza Hotel, Atlanta, Georgia, USA (2013)
34. Socher, R., Pennington, J., Huang, E.H., Ng, A.Y., Manning, C.D.: Semi-supervised recursive autoencoders for predicting sentiment-distributions. In: *Proceedings of the 2011 Conference on Empirical Methods in Natural Language Processing, EMNLP 2011, John McIntyre Conference Centre, Edinburgh, UK, A meeting of SIG-DAT, a Special Interest Group of the ACL*, pp. 151–161. ACL (2011)
35. Song, J., Yang, Y., Huang, Z., Shen, H.T., Hong, R.: Multiple feature hashing for real-time large scale near-duplicate video retrieval. In: *MM, ACM*, pp. 423–432 (2011)
36. Song, J., Yang, Y., Yang, Y., Huang, Z., Shen, H.T.: Inter-media hashing for large-scale retrieval from heterogeneous data sources. In: *SIGMOD Conference*, pp. 785–796 (2013)
37. Srivastava, N., Salakhutdinov, R.: Multimodal learning with deep boltzmann machines. In: *NIPS*, pp. 2231–2239 (2012)
38. Vincent, P., Larochelle, H., Bengio, Y., Manzagol, P.A.: Extracting and composing robust features with denoising autoencoders. In: *ICML*, pp. 1096–1103 (2008)
39. Wang, W., Ooi, B.C., Yang, X., Zhang, D., Zhuang, Y.: Effective multi-modal retrieval based on stacked auto-encoders. *PVLDB* **7**(8), 649–660(2014)
40. Weber, R., Schek, H.J., Blott, S.: A quantitative analysis and performance study for similarity-search methods in high-dimensional spaces. In: Gupta, A., Shmueli, O., Widom, J. (eds.) *Proceedings of 24rd International Conference on Very Large Data Bases*, New York, USA, pp. 194–205. Morgan Kaufmann (1998) (1998)
41. Weiss, Y., Torralba, A., Fergus, R.: Spectral hashing. In: Koller, D., Schuurmans, D., Bengio, Y., Bottou, L. (eds.) *Advances in Neural Information Processing Systems 21, Proceedings of the Twenty-Second Annual Conference on Neural Information Processing Systems, Vancouver, British Columbia, Canada*, pp. 1753–1760. Curran Associates, Inc., (2008)
42. Zhang, D., Agrawal, D., Chen, G., Tung, A.K.H.: Hashfile: an efficient index structure for multimedia data. In: *ICDE*, pp. 1103–1114. IEEE Computer Society, Hannover, Germany (2011)
43. Zhen, Y., Yeung, D.Y.: A probabilistic model for multimodal hash-function learning. In: Yang, Q., Agarwal, D., Pei, J. (eds.) *The 18th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, KDD '12, Beijing, China*, pp. 940–948. ACM (2012)
44. Zhu, X., Huang, Z., Shen, H.T., Zhao, X.: Linear cross-modal hashing for efficient multimedia search. In: *ACM Multimedia Conference, MM' 13, Barcelona, Spain*, pp. 143–152 (2013)
45. Zhuang, Y., Yang, Y., Wu, F.: Mining semantic correlation of heterogeneous multimedia data for cross-media retrieval. *IEEE Trans. Multimed.* **10**(2), 221–229 (2008)